

NAAN MUTHALVAN

IBM COLLABARATE

ARTIFICIAL INTELLIGENCE

PROJECT TITLE

MEASURE ENERGY CONSUMPTION

NAME : LOGESH E

DEPT & YEAR : CSE & III yr

REG.NO : 712221104004

COLLEGE : PARK COLLEGE OF ENGINEERING AND
TECHNOLOGY

BUILDING MODEL FOR ENERGY PREDICTION

- Hence, we are going to build our prediction model algorithm for measure energy consumption.
- It consist the process of loading the dataset, data transformation, feature importance, train/test split, Visualize feature to target relationship, modeling, Forecast on Test, Outlier Analysis, Reviewing : Train/Test Split, Feature Horizon, Lag Features, Train Using Cross Validation, Fold Analysis, Retraining on all Data and Predicting Future.
- Let we see the all process of building the model.

1.IMPORTING LIBRARIES AND DATA SET LOADING

CODE:

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns
import matplotlib.pyplot as plt
import xgboost as xgb
```

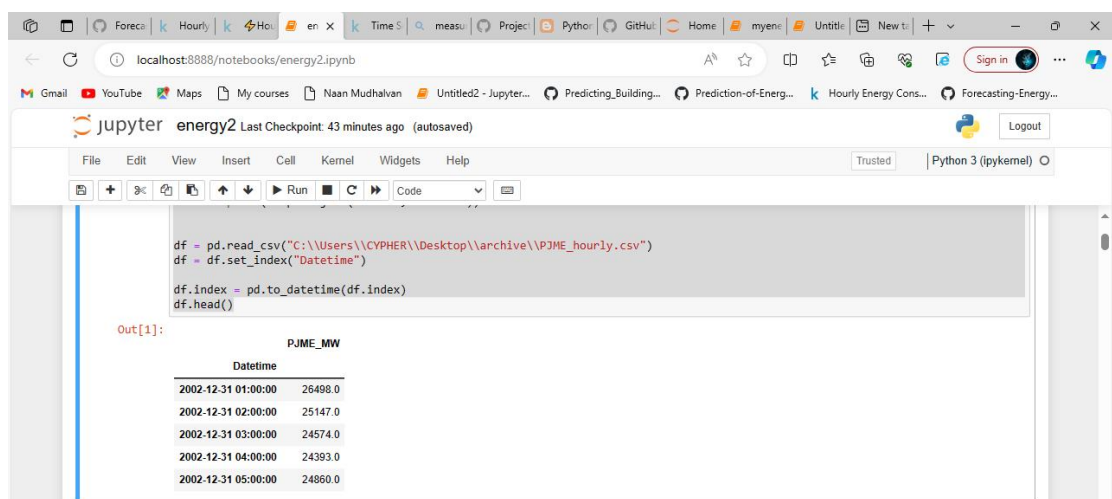
```
from sklearn.metrics import mean_squared_error
```

```
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
df =
pd.read_csv("C:\\Users\\CYPHER\\Desktop\\archive\\PJME_hourly.csv")
df = df.set_index("Datetime")
```

```
df.index = pd.to_datetime(df.index)
df.head()
```

O/P:



The screenshot shows a Jupyter Notebook window titled 'energy2'. The code cell contains the following Python code:

```
df = pd.read_csv("C:\\Users\\CYPHER\\Desktop\\archive\\PJME_hourly.csv")
df = df.set_index("Datetime")
df.index = pd.to_datetime(df.index)
df.head()
```

The output of the code is displayed below the code cell, showing the first five rows of the DataFrame:

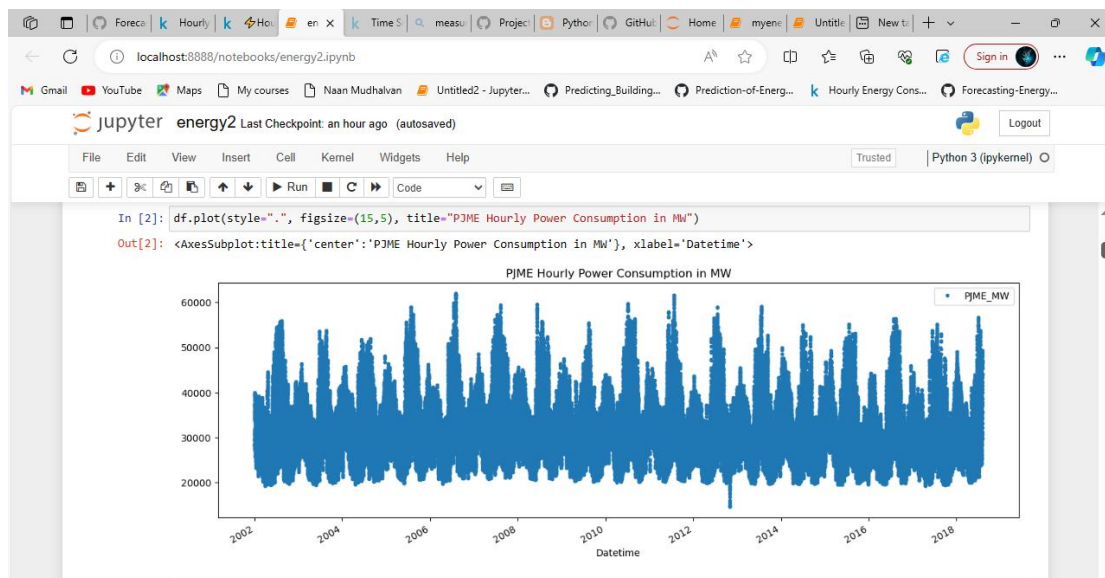
```
Out[1]:
```

PJME_MW	
Datetime	
2002-12-31 01:00:00	26490.0
2002-12-31 02:00:00	25147.0
2002-12-31 03:00:00	24574.0
2002-12-31 04:00:00	24393.0
2002-12-31 05:00:00	24860.0

2.TRAIN TEST SPLIT

CODE :

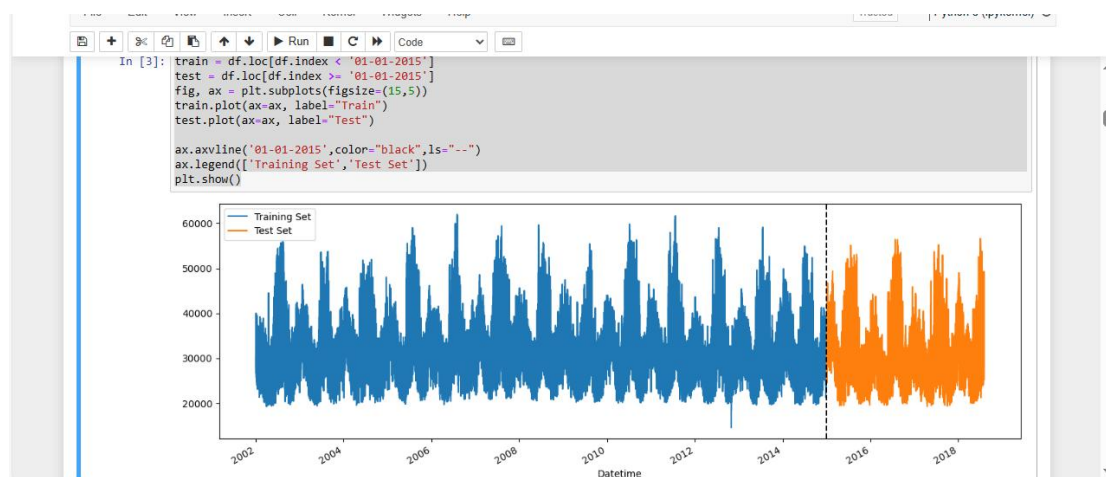
```
df.plot(style=".", figsize=(15,5), title="PJME Hourly Power Consumption in MW")
```



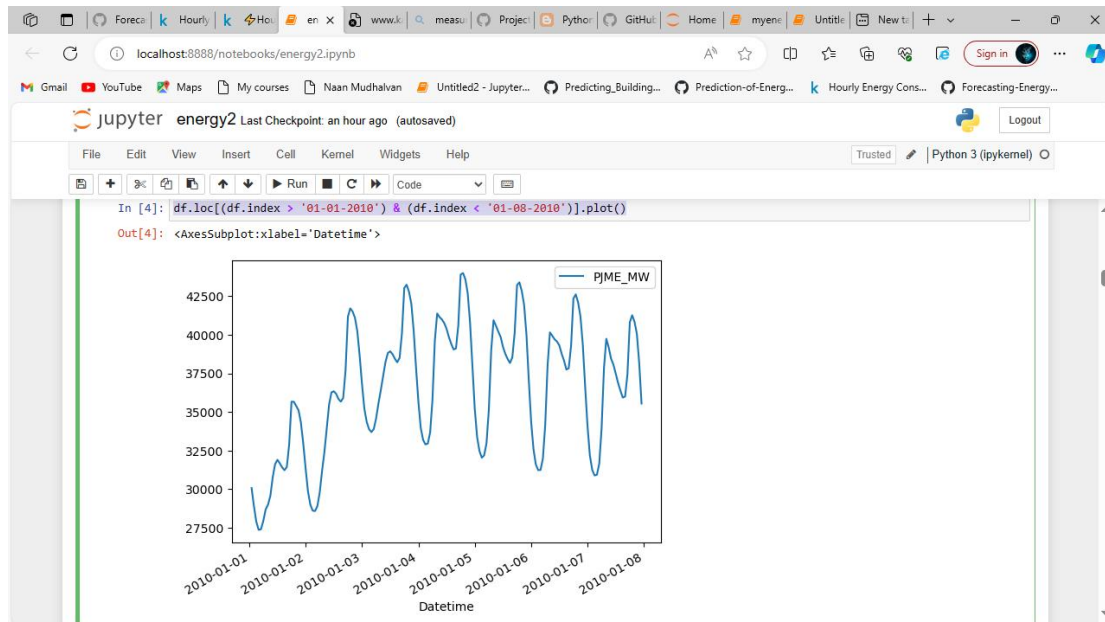
```
train = df.loc[df.index < '01-01-2015']
test = df.loc[df.index >= '01-01-2015']
fig, ax = plt.subplots(figsize=(15,5))
train.plot(ax=ax, label="Train")
test.plot(ax=ax, label="Test")
```

```
ax.axvline('01-01-2015',color="black",ls="--")
```

```
ax.legend(['Training Set','Test Set'])
plt.show()
```



```
df.loc[(df.index > '01-01-2010') & (df.index < '01-08-2010')].plot()
```



3. FEATURE CREATION

```
def create_time_series_features(dataframe):
```

```
    df = dataframe.copy()
    df['hour'] = df.index.hour
    df['dayofweek'] = df.index.dayofweek
    df['quarter'] = df.index.quarter
    df['month'] = df.index.month
    df['year'] = df.index.year
    df['dayofyear'] = df.index.dayofyear
```

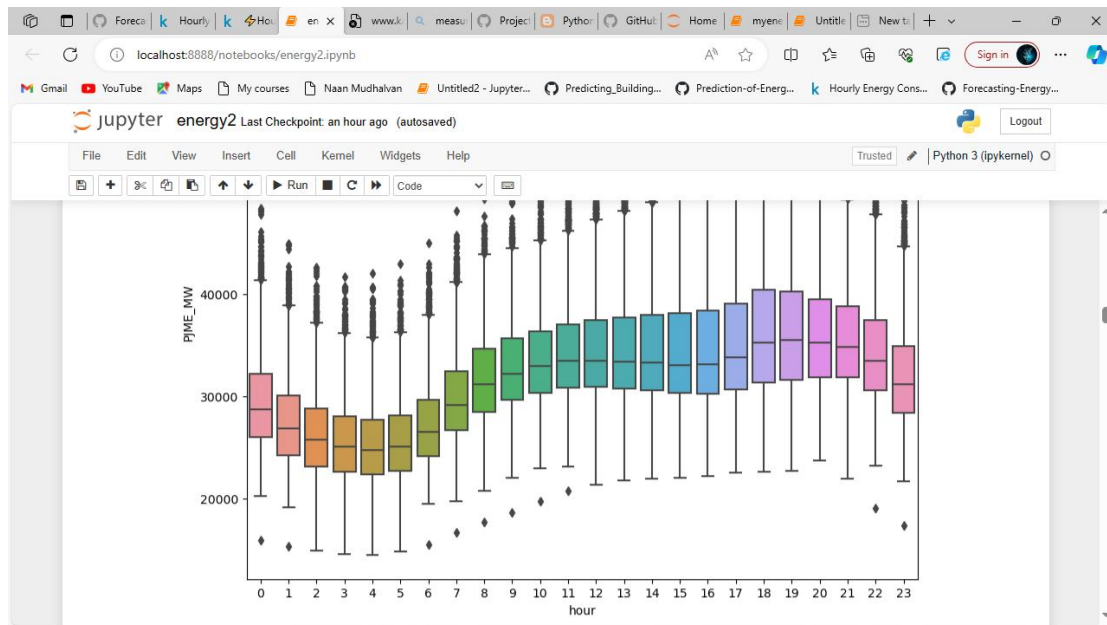
```
    return df
```

```
df = create_time_series_features(df)
```

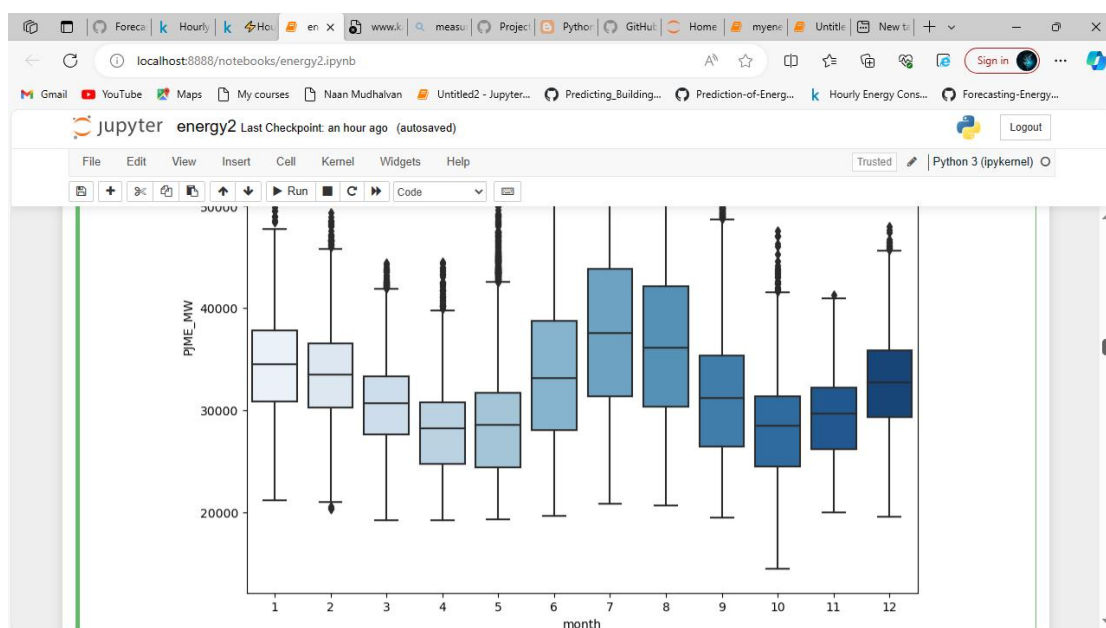
4. VISUALIZE FEATURE TO TARGET RELATIONSHIP

```
fig, ax = plt.subplots(figsize=(10,8))
sns.boxplot(data=df, x="hour", y="PJME_MW")
```

```
ax.set_title("MW By Hour")  
plt.show()
```

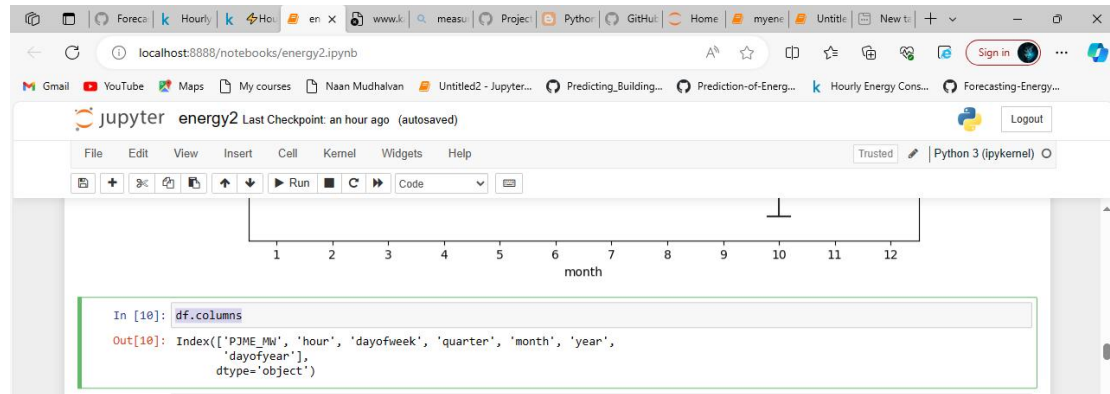


```
fig, ax = plt.subplots(figsize=(10,8))  
sns.boxplot(data=df, x="month", y="PJME_MW", palette="Blues")  
ax.set_title("MW By Month")  
plt.show()
```



5. MODELING

df.columns



```
FEATURES = ['hour', 'dayofweek', 'quarter', 'month', 'year',  
            'dayofyear']
```

```
OUTPUT = ['PJME_MW']
```

```
train = create_time_series_features(train)
```

```
test = create_time_series_features(test)
```

```
X_train = train[FEATURES]
```

```
y_train = train[OUTPUT]
```

```
X_test = test[FEATURES]
```

```
y_test = test[OUTPUT]
```

```
reg =
```

```
xg.XGBRegressor(n_estimators=1000,early_stopping_rounds=5  
0, learning_rate=0.01)
```

```
reg.fit(
```

```
    X_train,
```

```
    y_train,
```

```
    eval_set=[(X_train, y_train),(X_test, y_test)],
```

```
    verbose=100
```

```
)
```

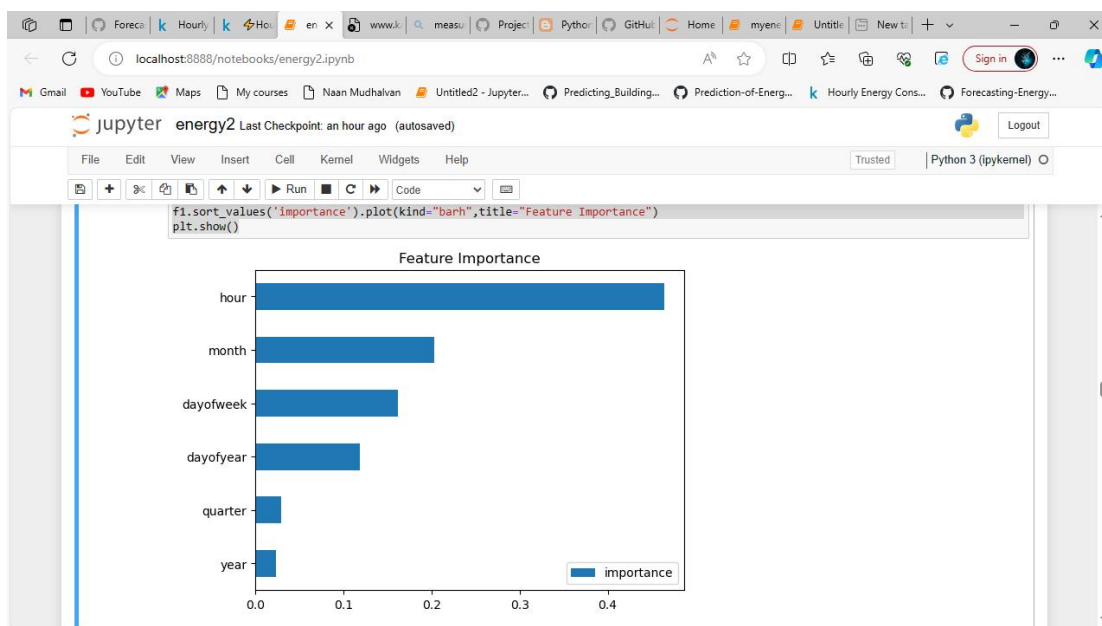
```
reg = xg.XGBRegressor(n_estimators=1000,early_stopping_rounds=50, learning_rate=0.01)
reg.fit(
    X_train,
    y_train,
    eval_set= [(X_train, y_train),(X_test, y_test)],
    verbose=100
)
```

[0]	validation_0-rmse:6407.35736	validation_1-rmse:6479.81619
[100]	validation_0-rmse:3911.97994	validation_1-rmse:4312.03224
[200]	validation_0-rmse:3244.38509	validation_1-rmse:3864.56545
[300]	validation_0-rmse:2996.08999	validation_1-rmse:3748.76687
[400]	validation_0-rmse:2830.28024	validation_1-rmse:3744.93348
[417]	validation_0-rmse:2801.66222	validation_1-rmse:3749.26889

```
Out[12]: XGBRegressor(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, device=None, early_stopping_rounds=50,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=0.01, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=None, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=1000, n_jobs=None,
    num_parallel_tree=None, random_state=None, ...)
```

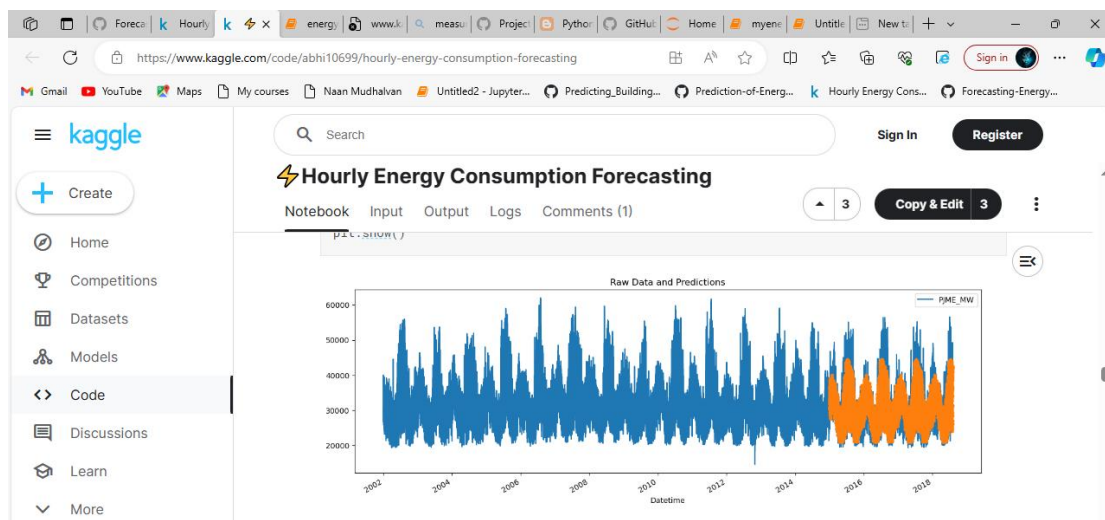
6.FEATURE IMPORTANCE

```
f1 = pd.DataFrame(data=reg.feature_importances_,
index=reg.feature_names_in_, columns=['importance'])
f1.sort_values('importance').plot(kind="barh",title="Feature
Importance")
plt.show()
```

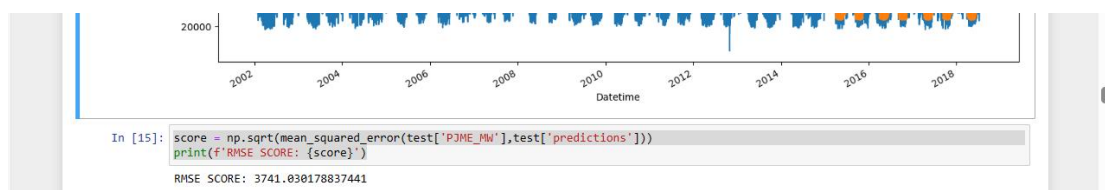


7.FEATURE FORECAST ON TEST

```
test['predictions'] = reg.predict(X_test)
df = df.merge(test[['predictions']], how='left', left_index=True,
right_index=True)
ax = df[['PJME_MW']].plot(figsize=(15,5))
df['predictions'].plot(ax=ax, style=".")
ax.set_title("Raw Data and Predictions")
plt.show()
```

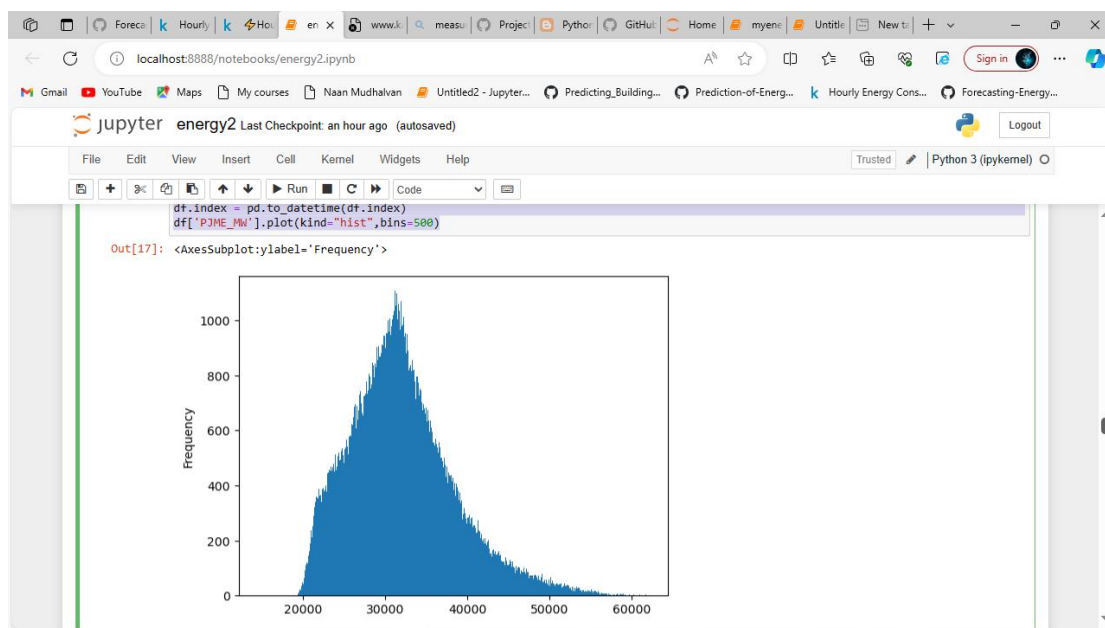


```
score =
np.sqrt(mean_squared_error(test['PJME_MW'],test['predictions']))
print(f'RMSE SCORE: {score}')
```



8. OUTLIER ANALYSIS

```
df =  
pd.read_csv("C:\\Users\\CYPHER\\Desktop\\archive\\PJME_hourly.csv")  
df = df.set_index("Datetime")  
  
df.index = pd.to_datetime(df.index)  
df['PJME_MW'].plot(kind="hist",bins=500)
```



```
df = df.query('PJME_MW > 19_000').copy()
```

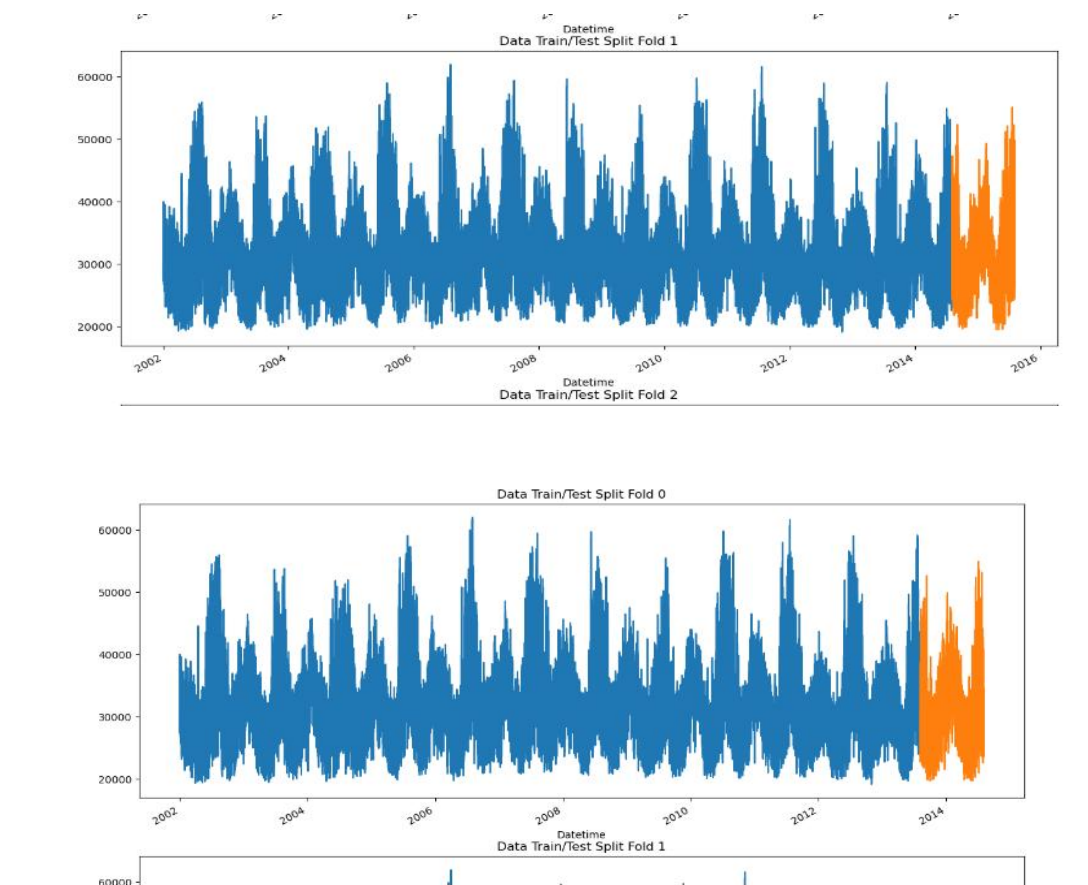
9. REVIEWING TRAIN AND TEST SPLIT

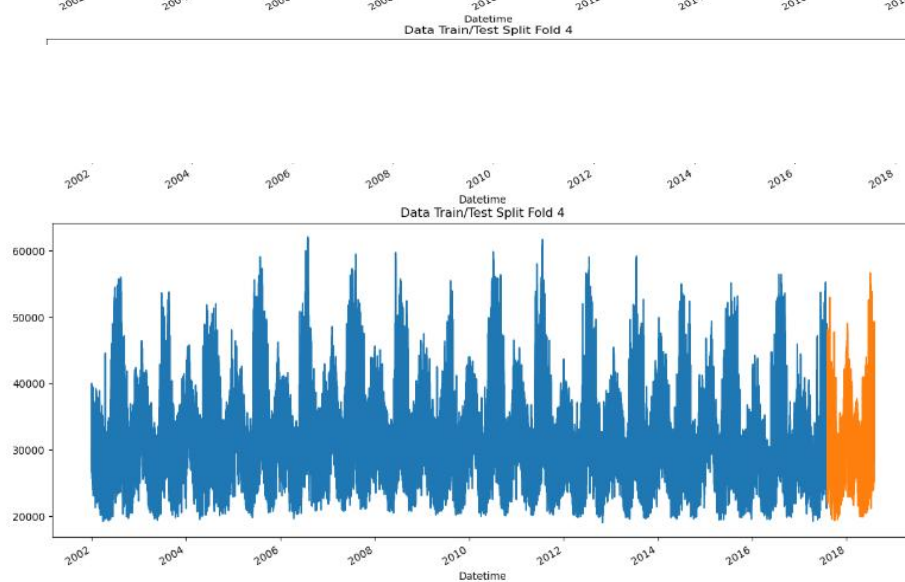
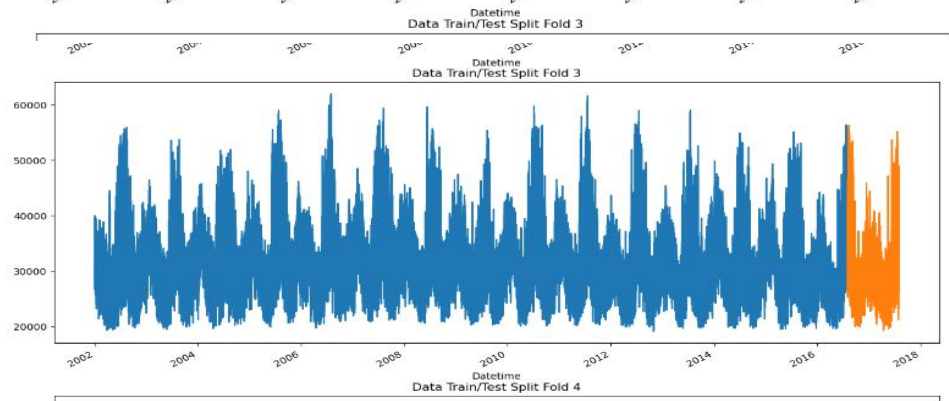
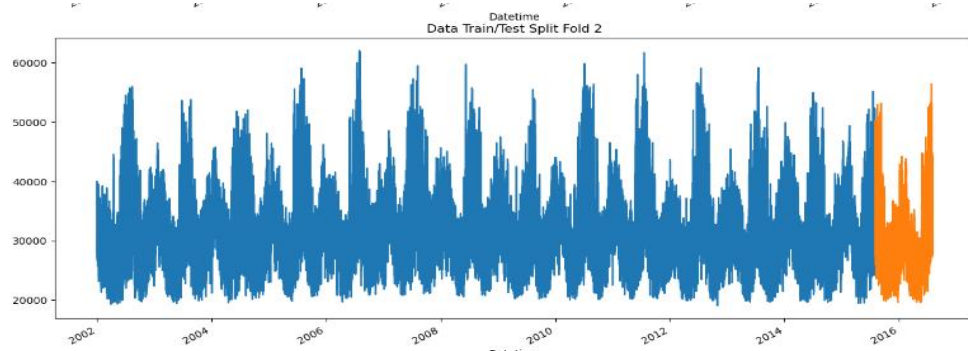
```
from sklearn.model_selection import TimeSeriesSplit  
tss = TimeSeriesSplit(n_splits=5, test_size=24*365*1,gap=24)  
df = df.sort_index()  
fig, axs = plt.subplots(5,1,figsize=(15,35))  
  
fold = 0  
for train_idx, val_idx in tss.split(df):
```

```
train = df.iloc[train_idx]
test = df.iloc[val_idx]
```

```
train['PJME_MW'].plot(
    ax=axes[fold],
    label="Training Set",
    title=f"Data Train/Test Split Fold {fold}"
)
test['PJME_MW'].plot(
    ax=axes[fold],
    label="Test Set",
)
```

```
fold += 1
```





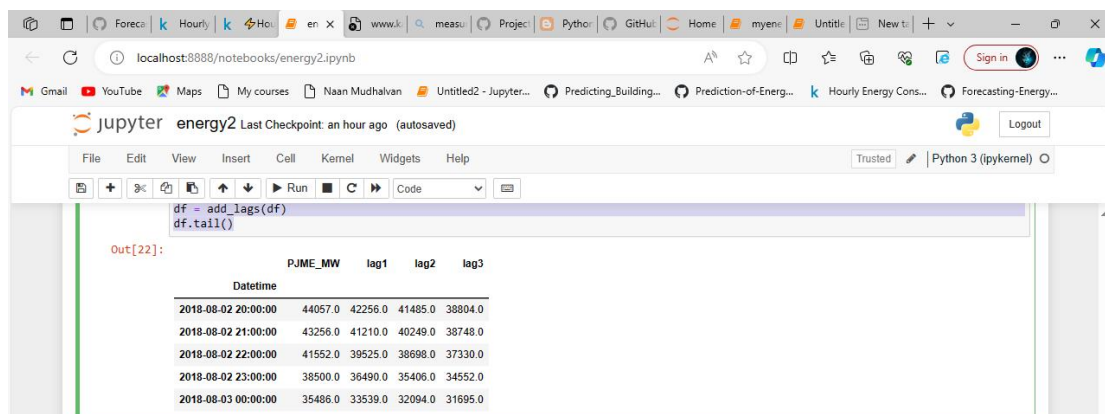
10. FEATURE HORIZON AND LAG FEATURES

```
df = create_time_series_features(df)
```

```
target_map = df['PJME_MW'].to_dict()
def add_lags(dframe):
    df = dframe.copy()
    df['lag1'] = (df.index - pd.Timedelta('364
days')).map(target_map)
    df['lag2'] = (df.index - pd.Timedelta('728
days')).map(target_map)
    df['lag3'] = (df.index - pd.Timedelta('1092
days')).map(target_map)

    return df
```

```
df = add_lags(df)
df.tail()
```



The screenshot shows a Jupyter Notebook window titled 'energy2'. The code cell contains the function definition and its application to the DataFrame. The output cell displays the last five rows of the DataFrame, showing the original PJME_MW values and their lags at 364, 728, and 1092 days prior.

	PJME_MW	lag1	lag2	lag3
2018-08-02 20:00:00	44057.0	42256.0	41485.0	38804.0
2018-08-02 21:00:00	43256.0	41210.0	40249.0	38748.0
2018-08-02 22:00:00	41552.0	39525.0	38698.0	37330.0
2018-08-02 23:00:00	38500.0	36490.0	35406.0	34552.0
2018-08-03 00:00:00	35486.0	33539.0	32094.0	31695.0

11. TRAIN USING CROSS VALIDATION

```
tss = TimeSeriesSplit(n_splits=5, test_size=24*365*1,gap=24)
df = df.sort_index()
df.columns
```

```
2018-08-02 23:00:00    38500.0    36490.0    35406.0    34552.0
2018-08-03 00:00:00    35486.0    33539.0    32094.0    31695.0

In [23]: tss = TimeSeriesSplit(n_splits=5, test_size=24*365*1,gap=24)
df = df.sort_index()
df.columns

Out[23]: Index(['PJME_MW', 'lag1', 'lag2', 'lag3'], dtype='object')
```

```
fold = 0
preds = []
scores = []
```

```
for train_idx, val_idx in tss.split(df):
    train = df.iloc[train_idx]
    test = df.iloc[val_idx]
```

```
train = create_time_series_features(train)
test = create_time_series_features(test)
```

```
FEATURES = ['hour', 'dayofweek', 'quarter', 'month', 'year',
'dayofyear',
    'lag1', 'lag2', 'lag3']
```

```
OUTPUT = 'PJME_MW'
```

```
OUTPUT = 'PJME_MW'
```

```
X_train = train[FEATURES]
```

```
y_train = train[OUTPUT]
```

```
X_test = test[FEATURES]
```

```
y_test = test[OUTPUT]
```

```
reg = xg.XGBRegressor(  
    base_score=0.5,  
    booster='gbtree',  
    n_estimators=1000,  
    early_stopping_rounds=50,  
    objective='reg:linear',  
    max_depth=3,  
    learning_rate=0.01  
)
```

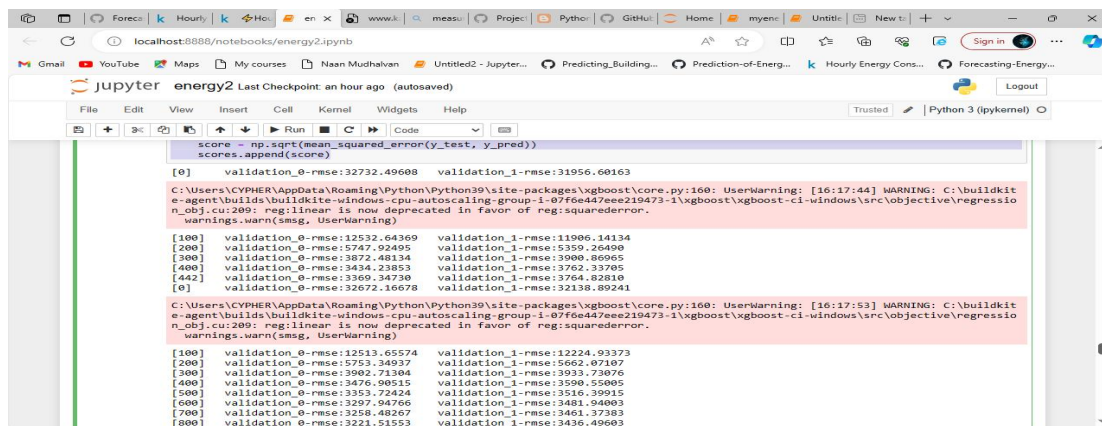
```
reg.fit(  
    X_train,  
    y_train,  
    eval_set=[(X_train, y_train),(X_test, y_test)],  
    verbose=100  
)
```

```
y_pred = reg.predict(X_test)
```

```
preds.append(y_pred)
```

```
score = np.sqrt(mean_squared_error(y_test, y_pred))
```

```
scores.append(score)
```



The screenshot shows a Jupyter Notebook window titled 'energy2'. The code cell contains the same Python code as the previous blocks. The output shows a list of RMSE scores for 1000 iterations, with a warning message from XGBoost.

```
score = np.sqrt(mean_squared_error(y_test, y_pred))  
scores.append(score)
```

```
[0] validation_0-rmse:32732.49608 validation_1-rmse:31956.60163  
[100] validation_0-rmse:12532.64369 validation_1-rmse:11906.14134  
[200] validation_0-rmse:5747.92495 validation_1-rmse:5359.26490  
[300] validation_0-rmse:3872.48134 validation_1-rmse:3900.86965  
[400] validation_0-rmse:3434.23853 validation_1-rmse:3762.33705  
[442] validation_0-rmse:3369.34730 validation_1-rmse:3764.82810  
[0] validation_0-rmse:32672.16678 validation_1-rmse:32138.89241  
[100] validation_0-rmse:12513.65574 validation_1-rmse:12224.93373  
[200] validation_0-rmse:5753.34937 validation_1-rmse:5662.07107  
[300] validation_0-rmse:3902.71304 validation_1-rmse:3933.73876  
[400] validation_0-rmse:3476.90515 validation_1-rmse:3590.55005  
[500] validation_0-rmse:3353.72424 validation_1-rmse:3516.39915  
[600] validation_0-rmse:3297.94766 validation_1-rmse:3481.94803  
[700] validation_0-rmse:3258.48267 validation_1-rmse:3461.37383  
[800] validation_0-rmse:3221.51553 validation_1-rmse:3436.49603
```

C:\Users\CYPHER\AppData\Roaming\Python\Python39\site-packages\xgboost\core.py:160: UserWarning: [16:17:44] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-1-07f6e447eee219473-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cc:209: reg:linear is now deprecated in favor of reg:squarederror.
warnings.warn(msg, UserWarning)

C:\Users\CYPHER\AppData\Roaming\Python\Python39\site-packages\xgboost\core.py:160: UserWarning: [16:17:53] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-1-07f6e447eee219473-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cc:209: reg:linear is now deprecated in favor of reg:squarederror.
warnings.warn(msg, UserWarning)


```
warnings.warn(msg, UserWarning)

[100] validation_0-rmse:12513.65574 validation_1-rmse:12224.93373
[200] validation_0-rmse:5753.34937 validation_1-rmse:5662.07187
[300] validation_0-rmse:3902.71304 validation_1-rmse:3933.73076
[400] validation_0-rmse:3476.90515 validation_1-rmse:3590.55005
[500] validation_0-rmse:3353.72424 validation_1-rmse:3516.39915
[600] validation_0-rmse:3297.94766 validation_1-rmse:3481.94003
[700] validation_0-rmse:3258.48267 validation_1-rmse:3461.37383
[800] validation_0-rmse:3221.51553 validation_1-rmse:3436.49603
[900] validation_0-rmse:3190.11480 validation_1-rmse:3428.88699
[999] validation_0-rmse:3166.16314 validation_1-rmse:3420.30469
[0] validation_0-rmse:32631.20370 validation_1-rmse:31073.29733

C:\Users\CYPHER\AppData\Roaming\Python\Python39\site-packages\xgboost\core.py:160: UserWarning: [16:18:09] WARNING: C:\buildkit
e-agent\builds\buildkite-windows-cpu-autoscaling-group-1-07f6e447ee219473-1\xgboost\xgboost-ci-windows\src\objective\regressio
n_obj.cu:209: reg:linear is now deprecated in favor of reg:squarederror.
  warnings.warn(msg, UserWarning)

[100] validation_0-rmse:12499.28425 validation_1-rmse:11136.70202
[200] validation_0-rmse:5750.81453 validation_1-rmse:4813.22087
[300] validation_0-rmse:3917.04200 validation_1-rmse:3553.46419
[400] validation_0-rmse:3494.55924 validation_1-rmse:3495.32356
[411] validation_0-rmse:3475.26636 validation_1-rmse:3503.65414
[0] validation_0-rmse:32528.44438 validation_1-rmse:31475.39670

C:\Users\CYPHER\AppData\Roaming\Python\Python39\site-packages\xgboost\core.py:160: UserWarning: [16:18:16] WARNING: C:\buildkit
e-agent\builds\buildkite-windows-cpu-autoscaling-group-1-07f6e447ee219473-1\xgboost\xgboost-ci-windows\src\objective\regressio
n_obj.cu:209: reg:linear is now deprecated in favor of reg:squarederror.
  warnings.warn(msg, UserWarning)
```

```
C:\Users\CYPHER\AppData\Roaming\Python\Python39\site-packages\xgboost\core.py:160: UserWarning: [16:18:16] WARNING: C:\buildkit
e-agent\builds\buildkite-windows-cpu-autoscaling-group-1-07f6e447ee219473-1\xgboost\xgboost-ci-windows\src\objective\regressio
n_obj.cu:209: reg:linear is now deprecated in favor of reg:squarederror.
  warnings.warn(msg, UserWarning)

[100] validation_0-rmse:12462.36581 validation_1-rmse:12020.28283
[200] validation_0-rmse:5738.57925 validation_1-rmse:5796.45874
[300] validation_0-rmse:3918.53218 validation_1-rmse:4388.39477
[400] validation_0-rmse:3501.24270 validation_1-rmse:4173.36380
[500] validation_0-rmse:3384.02490 validation_1-rmse:4119.70000
[600] validation_0-rmse:3325.50024 validation_1-rmse:4105.29234
[700] validation_0-rmse:3282.73755 validation_1-rmse:4091.57123
[800] validation_0-rmse:3250.37610 validation_1-rmse:4083.47152
[900] validation_0-rmse:3223.87614 validation_1-rmse:4081.83008
[999] validation_0-rmse:3199.82843 validation_1-rmse:4053.00975
[0] validation_0-rmse:32462.05557 validation_1-rmse:31463.90500

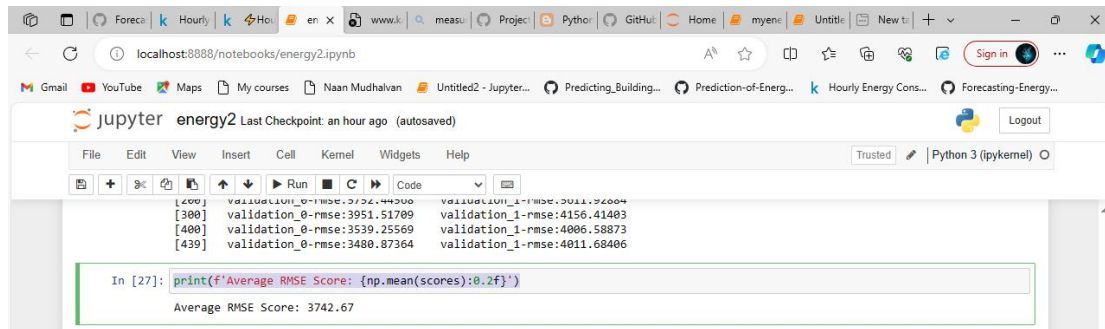
C:\Users\CYPHER\AppData\Roaming\Python\Python39\site-packages\xgboost\core.py:160: UserWarning: [16:18:34] WARNING: C:\buildkit
e-agent\builds\buildkite-windows-cpu-autoscaling-group-1-07f6e447ee219473-1\xgboost\xgboost-ci-windows\src\objective\regressio
n_obj.cu:209: reg:linear is now deprecated in favor of reg:squarederror.
  warnings.warn(msg, UserWarning)

[100] validation_0-rmse:12445.87740 validation_1-rmse:11963.42706
[200] validation_0-rmse:5752.44568 validation_1-rmse:5611.92884
[300] validation_0-rmse:3951.51709 validation_1-rmse:4156.41403
[400] validation_0-rmse:3539.25569 validation_1-rmse:4006.58873
[439] validation_0-rmse:3480.87364 validation_1-rmse:4011.68406

In [321]: print(f'Average RMSE Score: {rm_mean(rmse)}\n')
```


12.FOLD ANALYSIS & RETRAINING ON ALL DATA

```
print(f'Average RMSE Score: {np.mean(scores):0.2f}')
```



```
df = create_time_series_features(df)
```

```
FEATURES = ['hour', 'dayofweek', 'quarter', 'month', 'year', 'dayofyear',  
            'lag1', 'lag2', 'lag3']
```

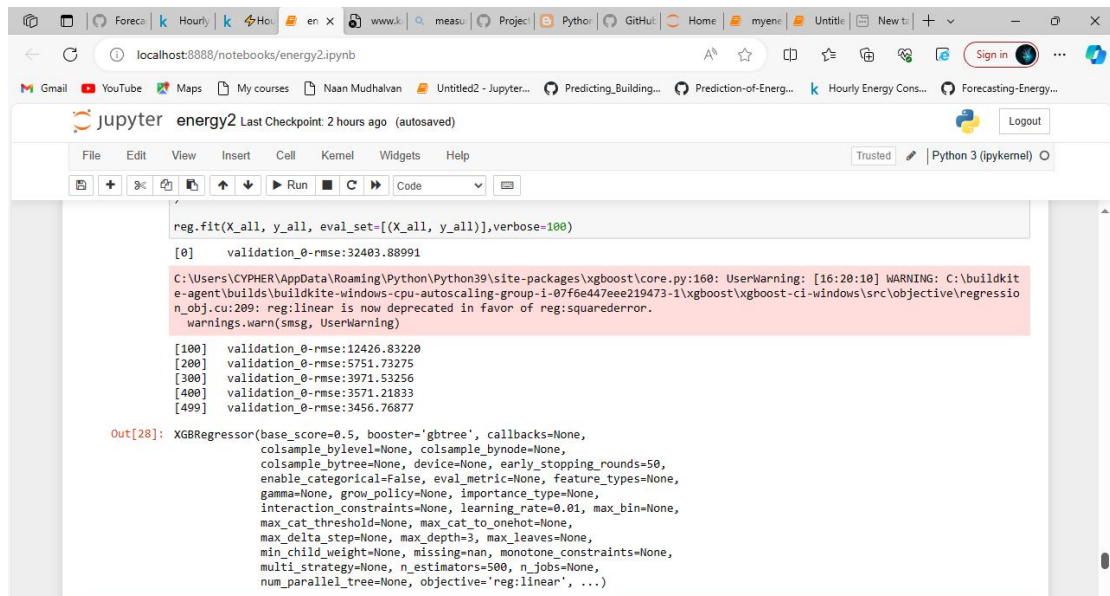
```
OUTPUT = 'PJME_MW'
```

```
X_all = df[FEATURES]
```

```
y_all = df[OUTPUT]
```

```
reg = xg.XGBRegressor(  
    base_score=0.5,  
    booster='gbtree',  
    n_estimators=500,  
    early_stopping_rounds=50,  
    objective='reg:linear',  
    max_depth=3,  
    learning_rate=0.01  
)
```

```
reg.fit(X_all, y_all, eval_set=[(X_all, y_all)], verbose=100)
```



The screenshot shows a Jupyter Notebook interface with a code cell containing the following Python code:

```
reg.fit(X_all, y_all, eval_set=[(X_all, y_all)], verbose=100)
```

The output of the code is a series of validation metrics:

```
[0]    validation_0-rmse:32403.88991
[100]  validation_0-rmse:12426.83220
[200]  validation_0-rmse:5751.73275
[300]  validation_0-rmse:3971.53256
[400]  validation_0-rmse:3571.21833
[499]  validation_0-rmse:3456.76877
```

A warning message is displayed in a red box:

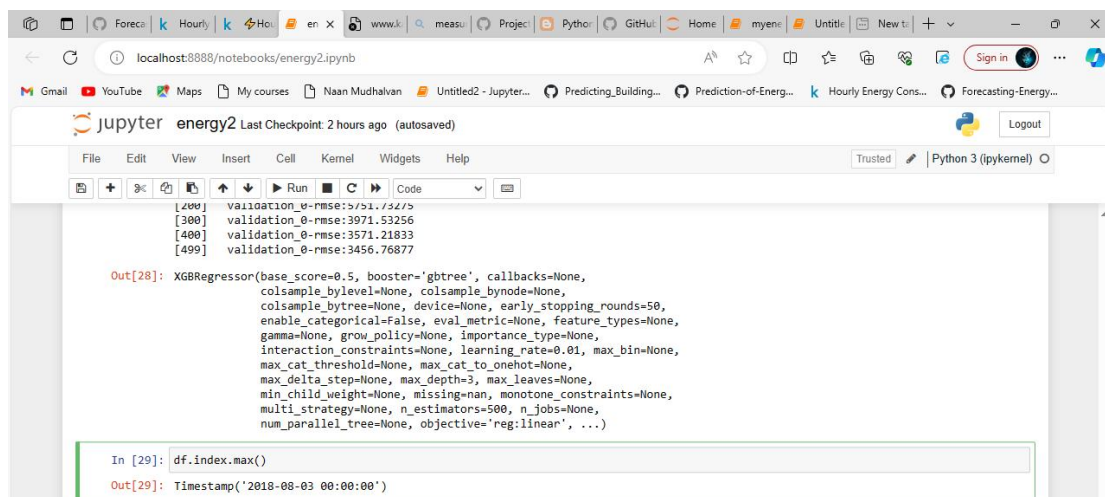
```
C:\Users\CYPHER\AppData\Roaming\Python\Python39\site-packages\xgboost\core.py:160: UserWarning: [16:20:10] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07f6e447eee219473-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:209: reg:linear is now deprecated in favor of reg:squarederror.
warnings.warn(msg, UserWarning)
```

The final output is the XGBoostRegressor object:

```
Out[28]: XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
colsample_bylevel=None, colsample_bynode=None,
colsample_bytree=None, device=None, early_stopping_rounds=50,
enable_categorical=False, eval_metric=None, feature_types=None,
gamma=None, grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=0.01, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=3, max_leaves=None,
min_child_weight=None, missing=None, monotone_constraints=None,
multi_strategy=None, n_estimators=500, n_jobs=None,
num_parallel_tree=None, objective='reg:linear', ...)
```

13.PREDICTING FUTURE

df.index.max()



The screenshot shows a Jupyter Notebook interface with a code cell containing the following Python code:

```
In [29]: df.index.max()
```

The output of the code is a timestamp:

```
Out[29]: Timestamp('2018-08-03 00:00:00')
```

```

future = pd.date_range('2018-08-03','2019-08-03',freq='1h')
future_df = pd.DataFrame(index=future)

future_df['isFuture'] = True

df['isFuture'] = False

df_and_future = pd.concat([df, future_df])

df_and_future = create_time_series_features(df_and_future)
df_and_future = add_lags(df_and_future)
future_w_features = df_and_future.query('isFuture').copy()
future_w_features['pred'] =
reg.predict(future_w_features[FEATURES])
future_w_features['pred'].plot(
    figsize=(10,5),
    ms=1,
    lw=1,
    title="Future Predictions"
)

plt.show()

```

